

Haofei Yu

haofei@cs.cmu.edu | haofeiyu.me | [github](https://github.com/haofeiyu) | [linkedin](https://www.linkedin.com/in/haofeiyu) | [twitter](https://twitter.com/haofeiyu)

Education

Carnegie Mellon University	2022/08 – 2024/05 (Expected)
Master of Science in Intelligent Information Systems, GPA 4.14 /4.33	Pittsburgh, PA, USA
<ul style="list-style-type: none">Teaching Assistant for 11-777 Multimodal Machine Learning (2023 Fall)A+ Coursework: Natural Language Processing, Advanced Topics in Multimodal Machine Learning, etc	
Zhejiang University	2018/09 – 2022/06
B.Eng in Computer Science and Technology (with Honors), GPA 3.96 /4.00	Hangzhou, China
<ul style="list-style-type: none">Rank: 7/134, Provincial Scholarship (top 5%)95+/100 Coursework: Discrete Mathematics and Application, Object-Oriented Programming, Computer Networks, etc	

Publications

[1] **Haofei Yu**, Uri Alon, Graham Neubig. *Racoon: Ranked Code Generation*. Under review.

[2] Xuhui Zhou*, Hao Zhu*, Leena Mathur, Ruohong Zhang, **Haofei Yu**, Zhengyang Qi, Louis-Philippe Morency, Yonatan Bisk, Daniel Fried, Graham Neubig, Maarten Sap. *SOTOPIA: Interactive Evaluation for Social Intelligence in Language Agents*. Under review.

[3] **Haofei Yu**, Paul Pu Liang, Ruslan Salakhutdinov, Louis-Philippe Morency. *Mixture of Multimodal Interaction Experts*. NeurIPS UniReps workshop.

[4] **Haofei Yu***, Cunxiang Wang*, Yue Zhang, Wei Bi. *TRAMS: Training-free Memory Selection for Long-range Language Modeling*. Findings of EMNLP 2023.

[5] Leonie Weissweiler*, Valentin Hofmann*, Anjali Kantharuban, Anna Cai, Ritam Dutt, Amey Hengle, Anubha Kabra, Atharva Kulkarni, Abhishek Vijayakumar, **Haofei Yu**, Hinrich Schuetze, Kemal Oflazer, David R Mortensen. *Counting the Bugs in ChatGPT’s Wugs: A Multilingual Investigation into the Morphological Capabilities of a Large Language Model*. EMNLP 2023.

[6] Cunxiang Wang*, **Haofei Yu***, Yue Zhang. *RFiD: Towards Rational Fusion-in-Decoder for Open-Domain Question Answering*. Findings of ACL 2023.

[7] Chiyu Song*, Hongliang He*, **Haofei Yu**, Leyang Cui, Pengfei Fang, Zhenzhong Lan. *Uni-Encoder: A Fast and Accurate Response Selection Paradigm for Generation-Based Dialogue Systems*. Findings of ACL 2023.

Work Experience

Apple - Siri & Information Intelligence	2023/05 – 2023/08
Machine Learning Intern	Seattle, WA, USA
<ul style="list-style-type: none">Enhanced user satisfaction for Siri’s long-range question answering and focused on complex, real-world user queries.Engineered a hierarchical prompting system to disambiguate and accurately respond to challenging questions.Chosen to present to Senior Director Robby Walker, showcasing intern project’s potential to augment Siri.	
Tencent - Tencent AI lab	2022/02 – 2022/08
Research Intern	Shenzhen, China
<ul style="list-style-type: none">Spearheaded a project to optimize inference efficiency in long-range language models, reducing perplexity on wiki-text-103 by 0.19 through an innovative memory selection method without training[4].	

Research Experience

Carnegie Mellon University - Language Technologies Institute	2022/09 – 2023/11
Research Assistant	Pittsburgh, PA, USA
<ul style="list-style-type: none">Worked with Prof. Graham Neubig for ranking-augmented code generation[1], Prof. Ruslan Salakhutdinov and Prof. Louis-Philippe Morency for social intelligence benchmark[2] and multimodal interaction[3].	
Westlake University - School of Engineering	2021/02 – 2022/02
Research Assistant	Hangzhou, China
<ul style="list-style-type: none">Worked with Prof. Zhenzhong Lan and Prof. Yue Zhang on open-domain natural language processing tasks including open-domain question answering[6] and open-domain dialogue generation[7].	